

Practicing and Extending Joins

DATA 202 21FA

Reminders

- Data Science Firepit Tonight!

Also:

- Quiz 5 closes tonight
- Homework 3 due Friday
- Discussion Forum replies due tomorrow

Preview: replication project

Gradescope trial

See Moodle.

Q&A

| Can't you rewrite any `right_join` as a flipped `left_join`?

Yes! It's mostly convenience (piped data comes in on *left*).

| Can we join more than 2 tables at once?

SQL? yes. `dplyr` makes you split it up though.

Q: Why did we need summarize()?

```
purchases_and_prices %>%  
  summarize(total_revenue = sum(price))
```

```
# A tibble: 1 × 1  
  total_revenue  
      <dbl>  
1          5.75
```

```
sum(price)
```

Error in eval(expr, envir, enclos): object 'price' not found

```
sum(purchases_and_prices$price)
```

```
[1] 5.75
```

Aside: vectors

So far we've always packaged vectors in data frames. But they can live on their own.

```
purchases_and_prices %>% select
```

```
# A tibble: 5 × 1
  price
  <dbl>
1     1
2    0.8
3    0.15
4    0.8
5     3
```

```
price_vec<- purchases_and_prices$price_vec
```

```
[1] 1.00 0.80 0.15 0.80 3.00
```

```
typeof(price_vec)
```

```
[1] "double"
```

Relational Data

- Bring together data that may never have been brought together before
- Values:
 - No one source has a complete view of the world
 - We need each other ("one body, many parts")
- Helps us get a more holistic picture
 - people aren't just their transactions; they also live somewhere

Example Applications

- Looking up abbreviations
- Connecting sales data to each customer's demographics
- Combining student data from Moodle, an online textbook, GitHub activity reports, ...

What others can you think of?

Code Together: Flight Delays

```
library(nycflights13)
```

```
flights %>%  
  drop_na(arr_delay) %>%  
  group_by(carrier) %>%  
  summarize(avg_delay = mean(arr_delay)) %>%  
  arrange(desc(avg_delay)) %>%  
  left_join(airlines, by = "carrier") %>%  
  select(name, avg_delay)
```

```
# A tibble: 16 × 2
```

	name	avg_delay
	<chr>	<dbl>
1	Frontier Airlines Inc.	21.9
2	AirTran Airways Corporation	20.1
3	ExpressJet Airlines Inc.	15.8
4	Mesa Airlines Inc.	15.6
5	SkyWest Airlines Inc.	11.9
6	Envoy Air	10.8
#	... with 10 more rows	

Multiple matches

`left_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4
		2	y5

Revenue by item?

purchases

customer_id	item
c1	bread
c1	milk
c1	banana
c2	milk
c2	toilet paper

prices

item	price
avocado	0.50
banana	0.15
bread	1.00
milk	0.80
toilet paper	3.00

Revenue by item?

For each item, look up all sales data.

```
prices %>%  
  left_join(purchases)
```

```
# A tibble: 6 × 3  
  item      price customer_id  
  <chr>    <dbl>    <chr>  
1 avocado  0.5      <NA>  
2 banana   0.15     c1  
3 bread    1        c1  
4 milk     0.8      c1  
5 milk     0.8      c2  
6 toilet paper 3        c2
```

Notice: multiple rows for each item. Where did each one come from?

Revenue by item?

```
prices %>%  
  left_join(purchases) %>%  
  group_by(item) %>%  
  summarize(revenue = sum(price))
```

```
# A tibble: 5 × 2  
  item      revenue  
  <chr>    <dbl>  
1 avocado    0.5  
2 banana    0.15  
3 bread      1  
4 milk      1.6  
5 toilet paper 3
```

...but nobody bought any avocados!

inner vs left

```
prices %>%  
  left_join(purchases)
```

```
# A tibble: 6 × 3  
  item           price customer_i  
  <chr>         <dbl> <chr>  
1 avocado       0.5   <NA>  
2 banana       0.15  c1  
3 bread         1     c1  
4 milk         0.8   c1  
5 milk         0.8   c2  
6 toilet paper  3     c2
```

```
prices %>%  
  inner_join(purchases)
```

```
# A tibble: 5 × 3  
  item           price customer_i  
  <chr>         <dbl> <chr>  
1 banana       0.15  c1  
2 bread         1     c1  
3 milk         0.8   c1  
4 milk         0.8   c2  
5 toilet paper  3     c2
```

Put it together

```
prices %>%  
  inner_join(purchases) %>%  
  group_by(item) %>%  
  summarize(revenue = sum(price))
```

```
# A tibble: 4 × 2  
  item          revenue  
  <chr>        <dbl>  
1 banana      0.15  
2 bread        1  
3 milk        1.6  
4 toilet paper 3
```


uh oh...

customer_id	item
c1	bread
c1	milk
c1	bananas
c2	milk
c2	toilet paper

item	price
AVOCADO	0.50
BANANA	0.15
BREAD	1.00
MILK	0.80
TOILET_PAPER	3.00

uh oh...

customer_id	item
c1	bread
c1	milk
c1	bananas
c2	milk
c2	toilet paper

item	price
AVOCADO	0.50
BANANA	0.15
BREAD	1.00
MILK	0.80
TOILET_PAPER	3.00

```
purchases %>%  
  left_join(prices)
```

```
# A tibble: 5 × 3  
  customer_id item      price  
  <chr>      <chr>    <dbl>  
1 c1        bread      NA  
2 c1        milk       NA  
3 c1        bananas    NA  
4 c2        milk       NA  
5 c2        toilet paper NA
```

Specifying keys

- Keys must match *exactly*
- Can join on multiple columns (first name **and** last name)
- Default join: columns with same names
- Specify what columns to use: `left_join(x, y, by = c("first_name", "last_name"))`

Make a new key that *does* match

```
prices_fixed <- prices %>%  
  mutate(item_norm = str_to_lower(item))  
purchases %>%  
  left_join(prices_fixed,  
            by = c("item" = "item_norm"))
```

```
# A tibble: 5 × 4  
  customer_id item      item.y price  
  <chr>      <chr>    <chr> <dbl>  
1 c1        bread    BREAD     1  
2 c1        milk     MILK     0.8  
3 c1        bananas <NA>      NA  
4 c2        milk     MILK     0.8  
5 c2        toilet paper <NA>      NA
```

But still some mismatches

```
purchases %>%  
  anti_join(  
    prices_fixed,  
    by = c("item" = "item_norm")
```

```
# A tibble: 2 × 2  
  customer_id item  
  <chr>      <chr>  
1 c1        bananas  
2 c2        toilet paper
```

```
prices_fixed %>%  
  anti_join(  
    purchases,  
    by = c("item_norm" = "item")
```

```
# A tibble: 3 × 3  
  item           price item_norm  
  <chr>         <dbl> <chr>  
1 AVOCADO         0.5 avocado  
2 BANANA          0.15 banana  
3 TOILET_PAPER    3 toilet_paper
```

Manual massaging: if_else

```
prices_fixed <- prices %>%  
  mutate(  
    item_norm = str_to_lower(item),  
    item_norm = if_else(  
      item_norm == "toilet_paper", "toilet paper",  
      item_norm)  
  )  
purchases %>%  
  left_join(prices_fixed, by = c("item" = "item_norm"))
```

```
# A tibble: 5 × 4  
  customer_id item      item.y      price  
  <chr>      <chr>      <chr>      <dbl>  
1 c1        bread      BREAD        1  
2 c1        milk      MILK         0.8  
3 c1        bananas   <NA>         NA  
4 c2        milk      MILK         0.8  
5 c2        toilet paper TOILET_PAPER 3
```

Manual massaging: case_when

```
prices_fixed <- prices %>%
  mutate(
    item_norm = str_to_lower(item),
    item_norm = case_when(
      item_norm == "toilet_paper" ~ "toilet paper",
      TRUE ~ item_norm)
  )
purchases %>%
  left_join(prices_fixed, by = c("item" = "item_norm"))
```

A tibble: 5 × 4

	customer_id	item	item.y	price
	<chr>	<chr>	<chr>	<dbl>
1	c1	bread	BREAD	1
2	c1	milk	MILK	0.8
3	c1	bananas	<NA>	NA
4	c2	milk	MILK	0.8
5	c2	toilet paper	TOILET_PAPER	3

Easy to add cases

```
prices_fixed <- prices %>%  
  mutate(  
    item_norm = str_to_lower(item),  
    item_norm = case_when(  
      item_norm == "toilet_paper" ~ "toilet paper",  
      item_norm == "banana" ~ "bananas",  
      TRUE ~ item_norm)  
  )  
purchases %>%  
  left_join(prices_fixed, by = c("item" = "item_norm"))
```

```
# A tibble: 5 × 4  
  customer_id item      item.y      price  
  <chr>      <chr>    <chr>    <dbl>  
1 c1        bread    BREAD      1  
2 c1        milk     MILK       0.8  
3 c1        bananas  BANANA     0.15  
4 c2        milk     MILK       0.8  
5 c2        toilet paper TOILET_PAPER 3
```


Aside: Could you use a `join` for this instead?

Think about this on your own.

case_when

if-elif version (Python):

```
if age < 0:  
    return "invalid"  
elif age < 18:  
    return "child"  
else:  
    return "adult"
```

case_when version:

```
age <- 18  
case_when(  
  age < 0 ~ "invalid",  
  age < 18 ~ "child",  
  TRUE ~ "adult"  
)
```

[1] "adult"

- first to True wins in both versions
- TRUE corresponds to else (the default)

case_when **vectorizes**

Like many R functions, it actually applies to all elements of a vector.

```
age <- c(-1, 0, 17, 18) # a vector
case_when(
  age < 0 ~ "invalid",
  age < 18 ~ "child",
  TRUE ~ "adult"
)
```

```
[1] "invalid" "child"   "child"   "adult"
```

case_when vs if_else

You can write the same thing either way. Which do you prefer?

if_else:

```
if_else(  
  age < 0, "invalid",  
  if_else(  
    age < 18, "child",  
    "other"))
```

```
[1] "invalid" "child"  "child"
```

case_when:

```
case_when(  
  age < 0 ~ "invalid",  
  age < 18 ~ "child",  
  TRUE ~ "adult"  
)
```

```
[1] "invalid" "child"  "child"
```

case_when in a data frame

```
people <- tribble(  
  ~name, ~age,  
  "Allen Linford", -1,  
  "Seb Dodds", 0,  
  "Charleen Lockwood", 17,  
  "Ridley Burgin", 18,  
)  
people %>% mutate(  
  adult = case_when(  
    age < 0 ~ "invalid",  
    age < 18 ~ "child",  
    TRUE ~ "adult"  
  )  
)
```

A tibble: 4 × 3

	name <chr>	age <dbl>	adult <chr>
1	Allen Linford	-1	invalid
2	Seb Dodds	0	child
3	Charleen Lockwood	17	child
4	Ridley Burgin	18	adult

The recoding pattern

```
population <- read_csv("../data/worldbank_sp_pop_totl.csv")
population %>% mutate(
  country = case_when(
    country == "United States" ~ "USA",
    iso3c == "GBR" ~ "UK", # LHS conditions may use
    TRUE ~ country # so can RHS
  )
) %>% filter(str_starts(country, "U")) # Just to see the results
```

More `case_when` tricks

See `?case_when` for how to:

- Deal with inconsistent data types
- Efficiently encode complicated conditionals
- Reuse `case_when` expressions by making a function

and more!